n = 83 field ops 2 (R1): local diagnostic

This is the local diagnostic that sized the Tier 1 design of `koblitz-n83-fieldops2-20261107`. It is not evidence: it ran on a shared cloud-agent VM, one pinned CPU (`taskset -c 2`), rustc 1.97, built from the frozen derivation (`make_fieldops2_source.py`) `[pcc:unverified]`. A is the field-ops row (`ic_online_fieldops_r770`), B the R1 producer (`ic_online_fieldops2_r770`), both built with `-C target-feature=+pclmulqdq`. Each round ran a fresh process of A and then of B on the M38 closure base: 4 cold rounds (seed 7) and 3 rank rounds (rank seed 1501) `[pcc:unverified]`.

stage (mean ms)	A	B	B / A
index build (cold process)	4240	4067	0.9592
scan self-test (cold process)	1171	1248	1.0656
online, one target (cold process)	861	820	0.9518
rank stage (rank process)	3531	3474	0.9840

On every round the index, layout, relation, per-attempt probe and root-call digests of A and B were identical `[pcc:unverified]`. The self-test does not run field products, so its B / A is noise on this VM and is kept as measured. The four B / A figures are the `modelled` factors that `power_check.py` reads (`inputs/local_profile.json`).